

test

December 10, 2022

1 Lab 04

1.1 Cross Validation

1.1.1 1-2

1) Reuse the notebook from Lab 3 for the wine data. Make sure to

* Reuse the same random seed throughout.

* Use nearest neighbors

2) With using KFold to produce the data splits, implement cross validation. Make sure to store the predictions on each test fold and print the `classification_report` after having looped over all folds.

```
[ ]: from sklearn.datasets import load_wine
      from sklearn.neighbors import KNeighborsClassifier
      from sklearn.preprocessing import StandardScaler
      from sklearn.metrics import classification_report
      from sklearn.model_selection import KFold
      import numpy as np

      x, y = load_wine(return_X_y = True) #split into features X and labels y
```

```
[ ]: kf = KFold(n_splits=3, random_state=None, shuffle=True)

      result_array = []
      y_test_report = []
      y_predict_report = []

      for train_index, test_index in kf.split(x):
          x_train, x_test = x[train_index], x[test_index]
          y_train, y_test = y[train_index], y[test_index]
          y_test_report.extend(y_test)
          scaler = StandardScaler(copy=True)
          xTrain_scaled = scaler.fit_transform(x_train, y_train)
          minDis = KNeighborsClassifier(n_neighbors=7)
          minDis.fit(xTrain_scaled, y_train)
```

```

xTest_scaled = scaler.transform(x_test)
y_predict_report.extend(minDis.predict(xTest_scaled))
result_array.append(minDis.score(xTest_scaled, y_test))

print('Average score: ', np.mean(result_array))

## print the test reports
print('The classification report:\n')
print(classification_report(y_test_report, y_predict_report))

```

1.1.2 3-4

3) Try with k=3 and k=10 folds.

4) In order to interpret the results (and fix possible issues), take a close look at the KFold visualization from the User Guide (not based on the wine data!):

```

[ ]: kf = KFold(n_splits=10, random_state=None, shuffle=True)

result_array = []
y_test_report = []
y_predict_report = []

for train_index, test_index in kf.split(x):
    x_train, x_test = x[train_index], x[test_index]
    y_train, y_test = y[train_index], y[test_index]
    y_test_report.extend(y_test)
    scaler = StandardScaler(copy=True)
    xTrain_scaled = scaler.fit_transform(x_train, y_train)
    minDis = KNeighborsClassifier(n_neighbors=7)
    minDis.fit(xTrain_scaled, y_train)
    xTest_scaled = scaler.transform(x_test)
    y_predict_report.extend(minDis.predict(xTest_scaled))
    result_array.append(minDis.score(xTest_scaled, y_test))

print('Average score: ', np.mean(result_array))
## print the test reports
print('The classification report:\n')
print(classification_report(y_test_report, y_predict_report))

```

Average score: 0.9663398692810456

The classification report:

	precision	recall	f1-score	support
0	0.95	1.00	0.98	59
1	1.00	0.92	0.96	71
2	0.94	1.00	0.97	48

accuracy			0.97	178
macro avg	0.96	0.97	0.97	178
weighted avg	0.97	0.97	0.97	178

Setting the shuffle parameter is very important since the classes are already ordered dataset

1.2 Grid Search

1.2.1 1-3

1) Implement Grid Search in combination with cross validation.

* Use the following parameters from the KNeighborsClassifier for the grid: n_neighbors and p . Select reasonable values for both.

* Implement a for loop to iterate over all combinations of the grid:

2) Run the Grid Search and print the classification report for each parameter combination.

3) Which parameter combination performs best?

```
[ ]: from sklearn.model_selection import ParameterGrid

n_neighbours = [2, 10]
p = [1, 2]

result_acb = {}

for n_nei in n_neighbours:
    for p_ in p:
        kf = KFold(n_splits=10, random_state=None, shuffle=True)
        result_array = []
        result_acb[str(n_nei) + " / " + str(p_)] = {}
        result_acb[str(n_nei) + " / " + str(p_)]["Y_TEST"] = []
        result_acb[str(n_nei) + " / " + str(p_)]["Y_PREDICT"] = []
        result_acb[str(n_nei) + " / " + str(p_)]["Y_Score"] = []

        for train_index, test_index in kf.split(x):
            x_train, x_test = x[train_index], x[test_index]
            y_train, y_test = y[train_index], y[test_index]
            result_acb[str(n_nei) + " / " + str(p_)]["Y_TEST"].extend(y_test)
            scaler = StandardScaler(copy=True)
            xTrain_scaled = scaler.fit_transform(x_train, y_train)
            minDis = KNeighborsClassifier(n_neighbors=n_nei, p=p_ )
```

```

minDis.fit(xTrain_scaled, y_train)
xTest_scaled = scaler.transform(x_test)
result_acb[str(n_nei) + " / " + str(p_)]["Y_PREDICT"].extend(minDis.
→predict(xTest_scaled))
result_acb[str(n_nei) + " / " + str(p_)]["Y_Score"].append(minDis.
→score(xTest_scaled, y_test))

## print the test reports
for parameters_in in result_acb.keys():
    print('Grid parameters:')
    print(parameters_in)
    print('Average score: ', np.mean(result_acb[parameters_in] ["Y_Score"]))
    print(classification_report(result_acb[parameters_in] ["Y_TEST"],
→result_acb[parameters_in] ["Y_PREDICT"]))
    □
→print('-----')

```

Grid parameters:

2 / 1

Average score: 0.9663398692810456

	precision	recall	f1-score	support
0	0.92	1.00	0.96	59
1	1.00	0.92	0.96	71
2	0.98	1.00	0.99	48
accuracy			0.97	178
macro avg	0.97	0.97	0.97	178
weighted avg	0.97	0.97	0.97	178

Grid parameters:

2 / 2

Average score: 0.9493464052287581

	precision	recall	f1-score	support
0	0.91	1.00	0.95	59
1	1.00	0.87	0.93	71
2	0.94	1.00	0.97	48
accuracy			0.95	178
macro avg	0.95	0.96	0.95	178
weighted avg	0.95	0.95	0.95	178

Grid parameters:

10 / 1

Average score: 0.9830065359477125

	precision	recall	f1-score	support
0	0.95	1.00	0.98	59
1	1.00	0.96	0.98	71
2	1.00	1.00	1.00	48
accuracy			0.98	178
macro avg	0.98	0.99	0.98	178
weighted avg	0.98	0.98	0.98	178

Grid parameters:

10 / 2

Average score: 0.9607843137254901

	precision	recall	f1-score	support
0	0.94	1.00	0.97	59
1	1.00	0.90	0.95	71
2	0.94	1.00	0.97	48
accuracy			0.96	178
macro avg	0.96	0.97	0.96	178
weighted avg	0.96	0.96	0.96	178

neighbour = 10 and manhattahn distance gets the best result

1.3 Combining Grid Search and Cross Validation

1.3.1 1 - 4

- 1) Carefully read the documentation of `GridSearchCV`, which combines the mechanisms of the grid search and the cross validation.
- 2) Reuse the `KNeighborsClassifier` and the `ParameterGrid` (check for correct naming).
- 3) Set the cross validation splitting strategy to `k=10` folds.
- 4) Evaluate the results using `GridSearchCV` 's built-in methods.

```
[ ]: from sklearn.model_selection import GridSearchCV
parameters = {"n_neighbors": [2, 10], "p": [1, 2]}
kn = KNeighborsClassifier()
clf = GridSearchCV(kn, parameters, cv = 10)

clf.fit(x, y)
```

```
print(clf.best_estimator_.score)
print("best score: ",clf.score(x,y))
```

```
<bound method ClassifierMixin.score of KNeighborsClassifier(n_neighbors=10,
p=1)>
best score: 0.8370786516853933
```

As we see the result is the same as we got in the previous task

1.3.2 5-6

5) Change the parameter scoring to use the F1 score for evaluation.

6) Find out how to store/access the best model parametrization.

```
[ ]: from sklearn.model_selection import GridSearchCV
parameters = {"n_neighbors": [2,10], "p": [1,2]}
kn = KNeighborsClassifier()
# f1_weighted because we do not have only 0 and 1 as values
clf = GridSearchCV(kn, parameters, cv = 10, scoring = "f1_weighted")

clf.fit(x,y)
estimator = clf.best_estimator_
print("best value for n_neighbors parameter: ", estimator.
      ↳get_params()['n_neighbors'])
print("best value for p parameter: ", estimator.get_params()["p"])
```

```
best value for n_neighbors parameter: 10
best value for p parameter: 1
```

1.4 Homework

Extend the grid with a parameter for switching the scaling of the data on/off. Then, for each test run made so far, enter the cross validation results in your table. Those values are more robust and reliable than those obtained from a single run.

```
[ ]: x, y = load_wine(return_X_y = True) #split into features X and labels y

parameters = {"n_neighbors": [2,10], "p": [1,2], "scale_with_mean": [False]}
kn = KNeighborsClassifier()
# f1_weighted because we do not have only 0 and 1 as values
clf = GridSearchCV(kn, parameters, cv = 10, scoring = "f1_weighted")

clf.fit(x,y)
estimator = clf.best_estimator_
print("best value for n_neighbors parameter: ", estimator.
      ↳get_params()['n_neighbors'])
print("best value for p parameter: ", estimator.get_params()["p"])
```

```

-----
ValueError                                Traceback (most recent call last)
/tmp/ipykernel_8436/2862377156.py in <module>
      6 clf = GridSearchCV(kn, parameters, cv = 10, scoring = "f1_weighted")
      7
----> 8 clf.fit(x,y)
      9 estimator = clf.best_estimator_
     10 print("best value for n_neighbors parameter: ", estimator.
      ↪ get_params()['n_neighbors'])

~/local/lib/python3.8/site-packages/sklearn/model_selection/_search.py in
      ↪ fit(self, X, y, groups, **fit_params)
     889             return results
     890
--> 891         self._run_search(evaluate_candidates)
     892
     893         # multimetric is determined here because in the case of a
      ↪ callable

~/local/lib/python3.8/site-packages/sklearn/model_selection/_search.py in
      ↪ _run_search(self, evaluate_candidates)
    1390     def _run_search(self, evaluate_candidates):
    1391         """Search all candidates in param_grid"""
-> 1392         evaluate_candidates(ParameterGrid(self.param_grid))
    1393
    1394

~/local/lib/python3.8/site-packages/sklearn/model_selection/_search.py in
      ↪ evaluate_candidates(candidate_params, cv, more_results)
     836         )
     837
--> 838         out = parallel(

     839             delayed(_fit_and_score)(
     840                 clone(base_estimator),

~/local/lib/python3.8/site-packages/joblib/parallel.py in __call__(self,
      ↪ iterable)
    1041         # remaining jobs.
    1042         self._iterating = False
-> 1043         if self.dispatch_one_batch(iterator):
    1044             self._iterating = self._original_iterator is not None
    1045

~/local/lib/python3.8/site-packages/joblib/parallel.py in
      ↪ dispatch_one_batch(self, iterator)
     859         return False

```

```

860             else:
--> 861                 self._dispatch(tasks)
862                 return True
863
~/.local/lib/python3.8/site-packages/joblib/parallel.py in _dispatch(self, batch)
777         with self._lock:
778             job_idx = len(self._jobs)
--> 779             job = self._backend.apply_async(batch, callback=cb)
780             # A job can complete so quickly than its callback is
781             # called before we get here, causing self._jobs to

~/.local/lib/python3.8/site-packages/joblib/_parallel_backends.py in
-> apply_async(self, func, callback)
206     def apply_async(self, func, callback=None):
207         """Schedule a func to be run"""
--> 208         result = ImmediateResult(func)
209         if callback:
210             callback(result)

~/.local/lib/python3.8/site-packages/joblib/_parallel_backends.py in
-> __init__(self, batch)
570         # Don't delay the application, to avoid keeping the input
571         # arguments in memory
--> 572         self.results = batch()
573
574     def get(self):

~/.local/lib/python3.8/site-packages/joblib/parallel.py in __call__(self)
260         # change the default number of processes to -1
261         with parallel_backend(self._backend, n_jobs=self._n_jobs):
--> 262             return [func(*args, **kwargs)

263                     for func, args, kwargs in self.items]
264

~/.local/lib/python3.8/site-packages/joblib/parallel.py in <listcomp>(.0)
260         # change the default number of processes to -1
261         with parallel_backend(self._backend, n_jobs=self._n_jobs):
--> 262             return [func(*args, **kwargs)

263                     for func, args, kwargs in self.items]
264

~/.local/lib/python3.8/site-packages/sklearn/utils/fixes.py in __call__(self,
-> *args, **kwargs)
209     def __call__(self, *args, **kwargs):
210         with config_context(**self.config):
--> 211             return self.function(*args, **kwargs)

```



```

212
213

~/local/lib/python3.8/site-packages/sklearn/model_selection/_validation.py in
↳ fit_and_score(estimator, X, y, scorer, train, test, verbose, parameters,
↳ fit_params, return_train_score, return_parameters, return_n_test_samples,
↳ return_times, return_estimator, split_progress, candidate_progress,
↳ error_score)
    667         cloned_parameters[k] = clone(v, safe=False)
    668
--> 669         estimator = estimator.set_params(**cloned_parameters)
    670
    671         start_time = time.time()

~/local/lib/python3.8/site-packages/sklearn/base.py in set_params(self,
↳ **params)
    238         key, delim, sub_key = key.partition("__")
    239         if key not in valid_params:
--> 240             raise ValueError(

    241                 "Invalid parameter %s for estimator %s. "
    242                 "Check the list of available parameters "

ValueError: Invalid parameter scale_with_mean for estimator
↳ KNeighborsClassifier(n_neighbors=2, p=1). Check the list of available
↳ parameters with `estimator.get_params().keys()`.

```